

Problem Statement:

For this model we will implement k-Means Clustering

```
In [1]: import pandas as pd  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [2]: df=pd.read_csv(r"C:\Users\my pc\Downloads\Online Retail.csv")
df
```

Out[2]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	01-12-2010 08:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	01-12-2010 08:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	01-12-2010 08:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	01-12-2010 08:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	01-12-2010 08:26	3.39	17850.0	United Kingdom
...
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	09-12-2011 12:50	0.85	12680.0	France
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	09-12-2011 12:50	2.10	12680.0	France
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	09-12-2011 12:50	4.15	12680.0	France
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	09-12-2011 12:50	4.15	12680.0	France
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	09-12-2011 12:50	4.95	12680.0	France

541909 rows × 8 columns

Data cleaning and Preprocessing

In [3]: `df.head()`

Out[3]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	01-12-2010 08:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	01-12-2010 08:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	01-12-2010 08:26	2.75	17850.0	United Kingdom
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	01-12-2010 08:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	01-12-2010 08:26	3.39	17850.0	United Kingdom

In [4]: `df.tail()`

Out[4]:

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	09-12-2011 12:50	0.85	12680.0	France
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	09-12-2011 12:50	2.10	12680.0	France
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	09-12-2011 12:50	4.15	12680.0	France
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	09-12-2011 12:50	4.15	12680.0	France
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	09-12-2011 12:50	4.95	12680.0	France

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   InvoiceNo        541909 non-null object
1   StockCode       541909 non-null object
2   Description     540455 non-null object
3   Quantity        541909 non-null int64
4   InvoiceDate      541909 non-null object
5   UnitPrice       541909 non-null float64
6   CustomerID      406829 non-null float64
7   Country         541909 non-null object
dtypes: float64(2), int64(1), object(5)
memory usage: 33.1+ MB
```

In [6]: `df.describe()`

Out[6]:

	Quantity	UnitPrice	CustomerID
count	541909.000000	541909.000000	406829.000000
mean	9.552250	4.611114	15287.690570
std	218.081158	96.759853	1713.600303
min	-80995.000000	-11062.060000	12346.000000
25%	1.000000	1.250000	13953.000000
50%	3.000000	2.080000	15152.000000
75%	10.000000	4.130000	16791.000000
max	80995.000000	38970.000000	18287.000000

In [7]: `df.shape`

Out[7]: (541909, 8)

```
In [8]: df.isna().any()
```

```
Out[8]: InvoiceNo      False
        StockCode     False
        Description    True
        Quantity      False
        InvoiceDate     False
        UnitPrice      False
        CustomerID     True
        Country        False
        dtype: bool
```

```
In [9]: df["InvoiceNo"].value_counts()
```

```
Out[9]: InvoiceNo
573585    1114
581219     749
581492     731
580729     721
558475     705
...
554023      1
554022      1
554021      1
554020      1
C558901      1
Name: count, Length: 25900, dtype: int64
```

```
In [10]: df=df.drop(["InvoiceNo","StockCode","Description","InvoiceDate"],axis=1)
```

```
In [11]: df
```

```
Out[11]:
```

	Quantity	UnitPrice	CustomerID	Country
0	6	2.55	17850.0	United Kingdom
1	6	3.39	17850.0	United Kingdom
2	8	2.75	17850.0	United Kingdom
3	6	3.39	17850.0	United Kingdom
4	6	3.39	17850.0	United Kingdom
...
541904	12	0.85	12680.0	France
541905	6	2.10	12680.0	France
541906	4	4.15	12680.0	France
541907	4	4.15	12680.0	France
541908	3	4.95	12680.0	France

541909 rows × 4 columns

```
In [12]: df.fillna(method="ffill",inplace=True) #filling the null values using forward fill method
```

```
In [13]: df.isna().any()
```

```
Out[13]: Quantity      False
UnitPrice      False
CustomerID      False
Country         False
dtype: bool
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: Quantity      0  
UnitPrice      0  
CustomerID      0  
Country      0  
dtype: int64
```

```
In [15]: df
```

```
Out[15]:
```

	Quantity	UnitPrice	CustomerID	Country
0	6	2.55	17850.0	United Kingdom
1	6	3.39	17850.0	United Kingdom
2	8	2.75	17850.0	United Kingdom
3	6	3.39	17850.0	United Kingdom
4	6	3.39	17850.0	United Kingdom
...
541904	12	0.85	12680.0	France
541905	6	2.10	12680.0	France
541906	4	4.15	12680.0	France
541907	4	4.15	12680.0	France
541908	3	4.95	12680.0	France

541909 rows × 4 columns


```
In [16]: df["Country"].value_counts()
```

```
Out[16]: Country
United Kingdom      495478
Germany             9495
France              8557
EIRE                 8196
Spain                2533
Netherlands          2371
Belgium              2069
Switzerland          2002
Portugal             1519
Australia            1259
Norway               1086
Italy                 803
Channel Islands      758
Finland              695
Cyprus                622
Sweden               462
Unspecified          446
Austria              401
Denmark              389
Japan                358
Poland               341
Israel               297
USA                  291
Hong Kong            288
Singapore            229
Iceland              182
Canada               151
Greece               146
Malta                127
United Arab Emirates  68
European Community   61
RSA                  58
Lebanon              45
Lithuania            35
Brazil               32
Czech Republic       30
Bahrain              19
Saudi Arabia         10
Name: count, dtype: int64
```



```
In [17]: plt.scatter(df["Quantity"],df["UnitPrice"])  
plt.xlabel("Quantity")  
plt.ylabel("UnitPrice")
```

```
Out[17]: Text(0, 0.5, 'UnitPrice')
```



```
In [18]: from sklearn.cluster import KMeans  
km=KMeans()  
km
```

```
Out[18]: 

▼ KMeans  
KMeans()


```

```
In [19]: y_predicted=km.fit_predict(df[["Quantity","UnitPrice"]])
y_predicted
```

C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

```
Out[19]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [20]: df["cluster"]=y_predicted
```


```
In [21]: df.head()
```

```
Out[21]:
```

	Quantity	UnitPrice	CustomerID	Country	cluster
0	6	2.55	17850.0	United Kingdom	0
1	6	3.39	17850.0	United Kingdom	0
2	8	2.75	17850.0	United Kingdom	0
3	6	3.39	17850.0	United Kingdom	0
4	6	3.39	17850.0	United Kingdom	0

```
In [22]: df1=df[df.cluster==0]
df2=df[df.cluster==1]
df3=df[df.cluster==2]
plt.scatter(df1["Quantity"],df1["UnitPrice"],color="red")
plt.scatter(df2["Quantity"],df2["UnitPrice"],color="green")
plt.scatter(df3["Quantity"],df3["UnitPrice"],color="blue")
plt.xlabel("Quantity")
plt.ylabel("UNITprice")
```

Out[22]: Text(0, 0.5, 'UNITprice')



```
In [23]: from sklearn.preprocessing import MinMaxScaler
```

```
In [24]: Scaler=MinMaxScaler()
```

```
In [25]: Scaler.fit(df[["Quantity"]])
```

```
Out[25]:
```

▼ MinMaxScaler

MinMaxScaler()

```
In [26]: df["Quantity"]=Scaler.transform(df[["Quantity"]])  
df.head()
```

```
Out[26]:
```

	Quantity	UnitPrice	CustomerID	Country	cluster
0	0.500037	2.55	17850.0	United Kingdom	0
1	0.500037	3.39	17850.0	United Kingdom	0
2	0.500049	2.75	17850.0	United Kingdom	0
3	0.500037	3.39	17850.0	United Kingdom	0
4	0.500037	3.39	17850.0	United Kingdom	0

```
In [27]: Scaler.fit(df[["UnitPrice"]])
```

```
Out[27]:
```

▼ MinMaxScaler

MinMaxScaler()

```
In [28]: df["UnitPrice"]=Scaler.transform(df[["UnitPrice"]])
df.head()
```

Out[28]:

	Quantity	UnitPrice	CustomerID	Country	cluster
0	0.500037	0.221150	17850.0	United Kingdom	0
1	0.500037	0.221167	17850.0	United Kingdom	0
2	0.500049	0.221154	17850.0	United Kingdom	0
3	0.500037	0.221167	17850.0	United Kingdom	0
4	0.500037	0.221167	17850.0	United Kingdom	0

```
In [29]: km=KMeans()
```

```
In [30]: y_predicted=km.fit_predict(df[["Quantity","UnitPrice"]])
y_predicted
```

C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

Out[30]: array([0, 0, 0, ..., 0, 0, 0])


```
In [31]: df["New Cluster"]=y_predicted
df.head()
```

Out[31]:

	Quantity	UnitPrice	CustomerID	Country	cluster	New Cluster
0	0.500037	0.221150	17850.0	United Kingdom	0	0
1	0.500037	0.221167	17850.0	United Kingdom	0	0
2	0.500049	0.221154	17850.0	United Kingdom	0	0
3	0.500037	0.221167	17850.0	United Kingdom	0	0
4	0.500037	0.221167	17850.0	United Kingdom	0	0

```
In [32]: df1=df[df["New Cluster"]==0]
df2=df[df["New Cluster"]==1]
df3=df[df["New Cluster"]==2]
plt.scatter(df1["Quantity"],df1["UnitPrice"],color="red")
plt.scatter(df2["Quantity"],df2["UnitPrice"],color="green")
plt.scatter(df3["Quantity"],df3["UnitPrice"],color="blue")
plt.xlabel("Age")
plt.ylabel("Income($)")
```

Out[32]: Text(0, 0.5, 'Income(\$)')




```
In [33]: km.cluster_centers_
```

```
Out[33]: array([[0.50005899, 0.22117195],  
               [0.49999657, 0.50519622],  
               [0.02092722, 0.22113061],  
               [0.49999588, 0.33389462],  
               [0.97907278, 0.22113061],  
               [0.49999383, 1.          ],  
               [0.50000617, 0.          ],  
               [0.50000432, 0.24394336]])
```

```
In [34]: df1=df[df["New Cluster"]==0]
df2=df[df["New Cluster"]==1]
df3=df[df["New Cluster"]==2]
plt.scatter(df1["Quantity"],df1["UnitPrice"],color="red")
plt.scatter(df2["Quantity"],df2["UnitPrice"],color="green")
plt.scatter(df3["Quantity"],df3["UnitPrice"],color="blue")
plt.scatter(km.cluster_centers_[0],km.cluster_centers_
            [1],color="orange",marker="+")
plt.xlabel("Age")
plt.ylabel("Income($)")
```

```
Out[34]: Text(0, 0.5, 'Income($)')
```





```
In [35]: k_rng=range(1,10)
sse=[]
for k in k_rng:
    km=KMeans(n_clusters=k)
    km.fit(df[["Quantity","UnitPrice"]])
    sse.append(km.inertia_)
```

```
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
C:\Users\my pc\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
```

```
In [36]: plt.plot(k_rng,sse)
plt.xlabel("k")
plt.ylabel("Sum of Squared Error")
```

```
Out[36]: Text(0, 0.5, 'Sum of Squared Error')
```



Conclusion:

Here we have clustered for the columns "Quantity","UnitPrice".We have founded clusters for two times, the new clusters formed are 0.Even though ,they are goodly clustered without any error

