

Problem Statement:

To perform an analytics report on 100 years of rainfall data

In [1]:

```
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

Data Collection:

In [2]:

```
train_df=pd.read_csv(r"C:\Users\raja\Downloads\rainfall in india.csv")
train_df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	O
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	38
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	19
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	18
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	22
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	26
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	11
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	14
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	7
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	16
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	16

4116 rows × 19 columns



In [3]:

```
test_df=pd.read_csv(r"C:\Users\raja\Downloads\district wise rainfall normal1.csv")
test_df
```

Out[3]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	A
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	27
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	42
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	46
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	42
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	71
...
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	52
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	63
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	35
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	59
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	21

641 rows × 19 columns



Data Cleaning and Preprocessing

In [4]:

```
train_df.head()
```

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	3
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	3
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	

In [5]:

```
test_df.head()
```

Out[5]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0

In [6]:

```
train_df.tail()
```

Out[6]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.3
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4

In [7]:

```
test_df.tail()
```

Out[7]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217

In [8]:

```
train_df.shape
```

Out[8]:

(4116, 19)

In [9]:

```
test_df.shape
```

Out[9]:

(641, 19)

In [10]:

```
train_df.size
```

Out[10]:

78204

In [11]:

```
test_df.size
```

Out[11]:

12179

In [12]:

```
train_df.describe()
```

Out[12]:

	YEAR	JAN	FEB	MAR	APR	MAY	
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.218659
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.745417
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.411100
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.359197
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.745417
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.127432
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.942432

In [13]:

```
test_df.describe()
```

Out[13]:

	JAN	FEB	MAR	APR	MAY	JUN	JUL
count	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000
mean	18.355070	20.984399	30.034789	45.543214	81.535101	196.007332	326.033697
std	21.082806	27.729596	45.451082	71.556279	111.960390	196.556284	221.364643
min	0.000000	0.000000	0.000000	0.000000	0.900000	3.800000	11.600000
25%	6.900000	7.000000	7.000000	5.000000	12.100000	68.800000	206.400000
50%	13.300000	12.300000	12.700000	15.100000	33.900000	131.900000	293.700000
75%	19.200000	24.100000	33.200000	48.300000	91.900000	226.600000	374.800000
max	144.500000	229.600000	367.900000	554.400000	733.700000	1476.200000	1820.900000

In [14]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   SUBDIVISION           4116 non-null   object
1   YEAR                  4116 non-null   int64
2   JAN                   4112 non-null   float64
3   FEB                   4113 non-null   float64
4   MAR                   4110 non-null   float64
5   APR                   4112 non-null   float64
6   MAY                   4113 non-null   float64
7   JUN                   4111 non-null   float64
8   JUL                   4109 non-null   float64
9   AUG                   4112 non-null   float64
10  SEP                   4110 non-null   float64
11  OCT                   4109 non-null   float64
12  NOV                   4105 non-null   float64
13  DEC                   4106 non-null   float64
14  ANNUAL                4090 non-null   float64
15  Jan-Feb               4110 non-null   float64
16  Mar-May               4107 non-null   float64
17  Jun-Sep               4106 non-null   float64
18  Oct-Dec               4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [15]:

test_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 641 entries, 0 to 640
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   STATE_UT_NAME         641 non-null    object
1   DISTRICT              641 non-null    object
2   JAN                   641 non-null    float64
3   FEB                   641 non-null    float64
4   MAR                   641 non-null    float64
5   APR                   641 non-null    float64
6   MAY                   641 non-null    float64
7   JUN                   641 non-null    float64
8   JUL                   641 non-null    float64
9   AUG                   641 non-null    float64
10  SEP                   641 non-null    float64
11  OCT                   641 non-null    float64
12  NOV                   641 non-null    float64
13  DEC                   641 non-null    float64
14  ANNUAL                641 non-null    float64
15  Jan-Feb              641 non-null    float64
16  Mar-May              641 non-null    float64
17  Jun-Sep              641 non-null    float64
18  Oct-Dec              641 non-null    float64
dtypes: float64(17), object(2)
memory usage: 95.3+ KB
```

In [16]:

train_df.isnull().sum()

Out[16]:

```
SUBDIVISION      0
YEAR              0
JAN               4
FEB               3
MAR               6
APR               4
MAY               3
JUN               5
JUL               7
AUG               4
SEP               6
OCT               7
NOV              11
DEC              10
ANNUAL           26
Jan-Feb           6
Mar-May           9
Jun-Sep          10
Oct-Dec          13
dtype: int64
```


In [17]:

```
test_df.isnull().sum()
```

Out[17]:

```
STATE_UT_NAME    0
DISTRICT         0
JAN              0
FEB              0
MAR              0
APR              0
MAY              0
JUN              0
JUL              0
AUG              0
SEP              0
OCT              0
NOV              0
DEC              0
ANNUAL           0
Jan-Feb          0
Mar-May          0
Jun-Sep          0
Oct-Dec          0
dtype: int64
```

In [18]:

```
train_df.dropna(inplace=True)
```

In [19]:

```
train_df.isnull().sum()
```

Out[19]:

```
SUBDIVISION      0
YEAR             0
JAN              0
FEB              0
MAR              0
APR              0
MAY              0
JUN              0
JUL              0
AUG              0
SEP              0
OCT              0
NOV              0
DEC              0
ANNUAL           0
Jan-Feb          0
Mar-May          0
Jun-Sep          0
Oct-Dec          0
dtype: int64
```

In [20]:

```
data=train_df[['JUL','JUN','AUG']]
```


Data Visualization

In [21]:

```
sns.pairplot(data)
```

Out[21]:

<seaborn.axisgrid.PairGrid at 0x2b57a757010>



In [22]:

```
test_df['STATE_UT_NAME'].value_counts()
```

Out[22]:

STATE_UT_NAME	
UTTAR PRADESH	71
MADHYA PRADESH	50
BIHAR	38
MAHARASHTRA	35
RAJASTHAN	33
TAMIL NADU	32
KARNATAKA	30
ORISSA	30
ASSAM	27
GUJARAT	26
JHARKHAND	24
ANDHRA PRADESH	23
JAMMU AND KASHMIR	22
HARYANA	21
PUNJAB	20
WEST BENGAL	19
CHATISGARH	18
ARUNACHAL PRADESH	16
KERALA	14
UTTARANCHAL	13
HIMACHAL	12
NAGALAND	11
MIZORAM	9
MANIPUR	9
DELHI	9
MEGHALAYA	7
SIKKIM	4
TRIPURA	4
PONDICHERRY	4
ANDAMAN And NICOBAR ISLANDS	3
GOA	2
DAMAN AND DIU	2
DADAR NAGAR HAVELI	1
CHANDIGARH	1
LAKSHADWEEP	1

Name: count, dtype: int64

In [23]:

```
test_df['DISTRICT'].value_counts()
```

Out[23]:

```
DISTRICT
BIJAPUR      2
BILASPUR     2
AURANGABAD   2
HAMIRPUR     2
NICOBAR      1
..
GONDA        1
GORAKHPUR    1
HARDOI       1
JAUNPUR      1
LAKSHADWEEP  1
Name: count, Length: 637, dtype: int64
```

Linear Regression

In [24]:

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

In [25]:


```
x=train_df[['JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN']]
y=train_df['ANNUAL']
```

In [26]:

```
sns.heatmap(x.corr())
```

Out[26]:

<Axes: >



In [27]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
```

In [28]:

```
regr = LinearRegression()
regr.fit(x_train,y_train)
print(regr.intercept_)
coeff_train_df=pd.DataFrame(regr.coef_,x.columns,columns=['coefficient'])
coeff_train_df
```

509.4300841827661

Out[28]:

	coefficient
JAN	1.783059
FEB	1.152417
MAR	0.612686
APR	0.704043
MAY	1.585119
JUN	2.915527

In [29]:

```
score=regr.score(x_test,y_test)
print(score)
```


0.8737805512901836

In [30]:

```
predictions=regr.predict(x_test)
plt.scatter(y_test,predictions)
```

Out[30]:

<matplotlib.collections.PathCollection at 0x2b57e7283a0>



In [31]:

```
x=np.array(train_df['JAN']).reshape(-1,1)
y=np.array(train_df['ANNUAL']).reshape(-1,1)
```

In [32]:


```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(x_train,y_train)
regr.fit(x_test,y_test)
```

Out[32]:

```
LinearRegression
LinearRegression()
```

In [33]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



Ridge and Lasso Regression

In [34]:

```
#Ridge Regression Model
from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

In [35]:

```
plt.figure(figsize=(10,10))
```

Out[35]:

<Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

In [36]:

```
features = train_df.columns[2:5]
target = train_df.columns[-1:]
#x and y values
x = train_df[features].values
y = train_df[target].values
#split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of x_train is (2863, 3)

The dimension of x_test is (1227, 3)

In [37]:

```
lr = LinearRegression()
#fit model
lr.fit(x_train ,y_train)
#predict
#prediction = lr.predict(x_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.014085252842472529

The test score for lr model is 0.013932760995991611

In [38]:

```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train,y_train)
test_score_ridge = ridgeReg.score(x_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.014084568755393767

The test score for ridge model is 0.013953801238720143

In [39]:

```
#Lasso Regression Model
lassoReg = Lasso(alpha=10)
lassoReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_lasso = lassoReg.score(x_train,y_train)
test_score_lasso = lassoReg.score(x_test,y_test)
print("\nLasso Model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso Model:

The train score for lasso model is 0.0035623805127747987


The test score for lasso model is 0.0044251341510098685

In [40]:

```
pd.Series(lassoReg.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[40]:

<Axes: >



In [41]:

```
#using the linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
#score
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.014082885214061869

0.013970064158112527

C:\Users\raja\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model_coordinate_descent.py:1568: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

In [42]:

```
#using the linear Cv model
from sklearn.linear_model import RidgeCV
#Ridge cross validation
ridge_cv = RidgeCV(alphas = [0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(x_train,y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(x_test,y_test)))
```

The train score for ridge model is 0.014084568755393767

The train score for ridge model is 0.013953801238729135

Elastic Net

In [43]:

```
from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
regr.score(x,y)
```

```
[ 0.40842497 -0.35386704  0.3507325 ]
[143.79960327]
```

Out[43]:

0.01442870782486183

In [44]:

```
y_pred_elastic = regr.predict(x_train)
```

In [45]:

```
mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
print("Mean squared Error on test set",mean_squared_error)
```

Mean squared Error on test set 27483.6226278673

Conclusion:

In the dataset we perform Linear Regression, Lasso, Ridge and Elastic Net. In Linear Regression we 87% accuracy .so this is the best model for this rainfall dataset